

KolmoX: Structural Preconditioning Architecture for Lossless Heterogeneous Data Compression

Transcending the Limits of Black-Box Entropy Coders via Bit-Exact Domain Decomposition

Author: KolmoX Open-Source Research Project

Date: August 2026

Specification: KMX2 Engine Spec (v1.3.0)

Correction notice - 2026-09-03. *The benchmark table and abstract of this document previously reported figures that could not be reproduced by the shipped code. Several described transforms that were not reachable from the domain router at all, and the numbers had been written into the document by a text-substitution script rather than measured. **All eleven rows have since been re-measured** against the current pipeline with bit-exact roundtrip assertions. Ten of the eleven use synthetic datasets and are marked †; only the Astrophysics FITS row rests on a real-world production capture. See the repository README for the same table and the methodology notes.*

Abstract

Conventional general-purpose compression algorithms (e.g., LZMA, Deflate, Zstandard) treat arbitrary inputs as untyped, one-dimensional scalar byte streams. While remarkably effective for textual data exhibiting repeated substring sequences, they systematically underperform on structured numerical, spatial, and floating-point domains—such as continuous trajectories, telemetry, multidimensional scientific floats, raster frames, and audio waveforms—due to high local entropy in lower-order bit planes.

In this paper, we present **KolmoX**, a two-stage hierarchical lossless compression architecture based on structural preconditioning. In the primary stage, a domain-aware engine identifies payload schemas and executes deterministic, bit-exact transforms: columnar channel demuxing, IEEE-754 *byte-plane slicing*, stride autocorrelation, and 2D spatial delta modeling. In the secondary stage, the separated streams are encapsulated into the multi-stream **KMX2** container and encoded via Finite State Entropy (FSE) using Zstandard.

Empirical evaluations across 11 distinct domains demonstrate consistent gains over baseline state-of-the-art compressors wherever the payload actually carries the structure a transform exploits. Measured end-to-end through the public pipeline API, KolmoX achieves **+64.10%** net reduction over Zstandard on CNC toolpath G-Code (15.84x ratio), **+63.12%** on a tessellated parametric CAD mesh (16.27x ratio), **+60.48%** on stereo PCM audio (2.59x ratio), and **+54.67%** on fixed-stride binary register frames, with decompression throughput reaching **~1094 MB/s** on the binary register path. These figures were obtained on synthetically generated datasets, disclosed as such in the table below. On the one real-world production corpus evaluated - a 117 MB JWST MIRI observation - the measured gain is **+16.31%** (1.47x → 1.76x).

Crucially, the architecture is designed to never lose to its own baseline: every domain transform competes against a plain Zstandard encoding of the same input, and the smaller of the two is what gets stored. On a high-entropy payload where no structure exists to exploit, the transform is discarded and the cost collapses to the 24-byte KMX2 header (measured: -0.46% on an x86 executable).

Comprehensive Benchmark Results

Every row was measured on 2026-09-03 through the public KolmoXPipeline.compress_bytes() / decompress_bytes() API - the path an actual user exercises - with a bit-exact roundtrip assertion on each. Figures obtained by invoking engine classes directly, outside the pipeline, are deliberately excluded: they omit the 24-byte KMX2 container header and bypass the adaptive competitive fallback, and were the source of the discrepancies this revision corrects. Every † row is reproducible by third parties with python benchmarks/benchmark_extended.py, which regenerates each dataset from a fixed seed and re-asserts every roundtrip. Compression ratios are deterministic; the throughput columns are single-run samples and vary by roughly ±15% between runs.

Data Domain	Transformation Pipeline	Baseline (Zstd L3)	KolmoX (KMX2)	Gain vs Zstd-3	Gain vs Zstd-19 ‡	Throughput (Comp / Decomp)
CNC G-Code (.gcode) †	Columnar Axis Separation	5.69x	15.84x	+64.10%	+57.33%	~24 MB/s / ~35 MB/s
Parametric 3D CAD Mesh (.obj) †	Prefix-Grouped Vertex Plane Slicing	6.00x	16.27x	+63.12%	+68.21%	~40 MB/s / ~80 MB/s
Audio PCM 16-bit (.wav) †	Stereo Mid/Side Decorrelation	1.02x	2.59x	+60.48%	+38.78%	~172 MB/s / ~284 MB/s
LiDAR XYZ Point Cloud † ¹	Columnar Coordinate Slicing	27.11x	65.79x	+58.79%	+23.37%	~80 MB/s / ~106 MB/s
Temporal Video Sequence (.kmxvraw) †	Frame Arithmetic Delta (SIMD)	1.00x	2.31x	+56.74%	+60.22%	~148 MB/s / ~168 MB/s
Binary Register Packets (.bin) †	Stride Autocorr + Demux	1.86x	4.10x	+54.67%	+44.38%	~158 MB/s / ~1094 MB/s
Industrial Telemetry (.csv) †	Shape-Grouped Columnar Demux	4.59x	8.23x	+44.24%	+35.59%	~47 MB/s / ~98 MB/s
2D Natural Sensor Raster (.bmp) †	2D Spatial Delta	1.07x	1.80x	+40.38%	+33.59%	~37 MB/s / ~2.4 MB/s
Temporal Video (real gameplay) ³	Frame Arithmetic Delta (SIMD)	1.70x	2.79x	+38.87%	not measured	~87 MB/s / ~59 MB/s
Dense Float32 Buffer (250k values) †	IEEE-754 Byte-Plane Slicing	1.29x	1.89x	+31.86%	+20.47%	~236 MB/s / ~642 MB/s
Astrophysics FITS (JWST) - real data	IEEE-754 Byte-Plane Slicing (auto-routed)	1.47x	1.76x	+16.31%	+20.64%	~213 MB/s / ~714 MB/s
x86 Binary Executable (.exe) † ²	Adaptive Fallback (BCJ/Zstd)	7.59x	7.56x	-0.46%	-0.46%	~13 MB/s / ~879 MB/s

† Synthetic dataset, generated programmatically. Only the Astrophysics FITS row uses a real-world production capture: a 117 MB NASA/STScI JWST MIRI observation of the Carina Nebula, compressed whole-file through the automatic domain router. Compressing only the extracted pixel plane instead yields 1.26x → 1.61x (+21.41%).

¹ The LiDAR dataset is a deterministic arithmetic ramp, which is why even the plain Zstd baseline reaches 27.11x. It demonstrates that the transform functions, but it is **not** representative of real scanner output and should be read as a mechanism check rather than a performance claim.

² A 40 KB synthetic instruction stream; at that size the throughput timings are dominated by measurement noise.

‡ Obtained with benchmarks/benchmark_extended.py --strong-baselines. The comparison is like-for-like: level 19 is applied to both sides, so KolmoX's internal compressor and its adaptive fallback both operate at 19 against a Zstd-19 baseline. This is what makes the figure meaningful - measuring KolmoX at level 3 against Zstd-19 would only demonstrate that level 19 compresses better, and would say nothing about the value of the structural preconditioning itself. The real-gameplay row comes from a private capture outside the reproducible script and was not re-run at level 19. The same flag also prints a Gzip-9 column, which is a reference point rather than a measurement of structural value: the pipeline uses Zstd internally and its backend is not configurable, so a Gzip comparison conflates the preconditioning with Zstd-vs-Gzip. Only the Zstd-19 column isolates the former.

Computational cost: level 19 is roughly two orders of magnitude slower to compress (Zstd-3 reaches 179-2802 MB/s on these datasets, Zstd-19 only 3-13 MB/s). KolmoX at level 19 is slower still, 1.4-10.1 MB/s, because the adaptive fallback compresses

twice - once for the baseline it must beat, once for the candidate. Negligible at level 3; at level 19 it doubles the most expensive operation in the pipeline.

Reading the two gain columns. The preconditioning retains its value against a far more aggressive entropy coder: no domain drops to zero or turns negative, and the transform is still selected wherever it was at level 3. How much value, however, depends heavily on the domain - from +68.21% on the CAD mesh down to +23.37% on LiDAR. Three domains improve at level 19, which a single column would have concealed. And the row that loses the most, LiDAR, is precisely the one whose dataset was artificially favourable to begin with: a deterministic arithmetic ramp that Zstd-19 discovers unaided. That is independent corroboration of the caveat in note ¹ rather than a coincidence.

³ Real data: 60 consecutive frames of 5120x1440 gameplay footage at full resolution, no downscaling. It sits well below the synthetic video row because the content is adversarial for a temporal transform - 61% of bytes change between consecutive frames, up to 89% in the busiest pairs. The transform still wins the competitive check. Read the two video rows together: the synthetic figure is what low-motion content gives, not what video gives in general.

Both video rows improved when the temporal transform moved from an XOR delta to an arithmetic delta mod 256. XOR is a poor fit for continuous data: two adjacent values straddling a high bit yield a large residual ($127 \wedge 128 = 255$) where subtraction yields 1. On the real footage this took the gain from +25.14% to +38.87%, an output 18.34% smaller. The payload is versioned by its magic - KMXV2 is written, KMXV1 from earlier releases still decodes bit-exact.

Note: Domain gains depend on the input actually carrying the structure the transform exploits. On a high-entropy payload - a random-walk mesh with six decimals of noise per coordinate, or the x86 stream above - the transform yields nothing, the adaptive competitive fallback discards it, and the stored result is the plain Zstd baseline plus the 24-byte KMX2 header. The raster decompression figure (~2.4 MB/s) reflects a pure-Python pixel loop whose sequential dependency has not yet been ported to the C extension; it is the current honest number, not a target.

Real-World Data Benchmark

The table above is measured on synthetic datasets generated from fixed seeds.

This second table is measured on real datasets drawn from public archives with declared licenses, reproducible with `python benchmarks/download_datasets.py` followed by `python benchmarks/benchmark_real.py` (~52 MB of download). It is kept separate rather than merged into a further column, because the underlying data differs and combining the two would obscure both.

The pattern it reveals is worth stating plainly. Where the parser recognises the real data the gain holds - FITS retains +14.83% against +16.31% on its synthetic counterpart, and a well-formed CSV retains +24.85%. Where it does not, the gain collapses and the adaptive competitive fallback correctly stores a plain Zstd encoding instead.

Two domains collapsed when this benchmark was first run, and both proved to be defects of dialect recognition rather than limits of structural preconditioning.

Neither was repaired by strengthening a transform; in both cases the existing transform merely began to see the file's actual shape. G-code moved from -0.06% to +21.93% once attached line numbers and modal commands were handled, and the semicolon-delimited CSV moved from +11.51% to +26.29% once the delimiter was inferred rather than assumed. The footnotes qualify the former: it improves for

a reason that generalises less readily than the figure implies.

Dataset	Source	License	Size	Baseline (Zstd L3)	KolmoX	Gain	Outcome
Astrophysics FITS (JWST SMACS 0723)	MAST / STScI	Public domain	35.0 MB	1.76x	2.07x	+14.83%	transform used
Industrial Telemetry (clean CSV)	UCI ML Repository	CC BY 4.0	11.4 MB	6.02x	8.01x	+24.85%	transform used
Industrial Telemetry (semicolon CSV)	UCI ML Repository	CC BY 4.0	0.8 MB	3.04x	4.13x	+26.29%	transform used (a)
CNC G-Code (LinuxCNC)	LinuxCNC	GPL-2.0 (b)	0.2 MB	4.79x	6.13x	+21.93%	transform used (c)
CAD Mesh (binary STL, WVS)	Zenodo 5034614	CC0	3.1 MB	1.86x	1.86x	-0.00%	fallback (d)
CAD Mesh (binary STL, Ambulacral)	Zenodo 5034614	CC0	0.7 MB	2.57x	2.56x	-0.01%	fallback (d)

(a) Semicolon-delimited with commas as decimal separators, as European locales write them. The demux previously assumed commas, splitting the decimals and producing ragged rows, and this row read +11.51%. The delimiter is now inferred from the data, recovering the full 14.78 points the diagnosis had predicted. The improvement is not a stronger transform but the same transform correctly recognising the file's dialect - the second of the two collapses this benchmark exposed, both now closed. The clean CSV and the synthetic dataset are unchanged: detection selects the comma for them and the code path is byte-identical to before.

(b) Downloaded for local benchmarking only; the file carries an internal copyright notice and is not redistributed with this repository.

(c) **This gain derives from line numbers, not coordinates, and should not be read as evidence that G-code preconditioning transfers to real files.** On this file the engine extracts X=0, Y=0, Z=0, N=4691. The 4,691 extracted values are RS-274/NGC line numbers - metadata - and collapsing them out of the template is the entire source of the +21.93%. The coordinates, 4,684 values constituting the actual payload, remain untouched: this file expresses them as parametric expressions (Y[#<yscale>*53.293]) rather than literals, and the engine deliberately does not model that dialect's expression syntax. Measurable headroom consequently remains. Prior to the fix this row read -0.06% with the transform discarded, because the engine matched only lines beginning with G1 while 4,690 of 4,691 lines here attach the line number with no separating space and omit the G-word entirely (modal G-code). Both cases are now handled; the earlier and incorrect diagnosis is recorded in TODO.md.

(d) The two obvious remedies were measured and rejected: transposing at the 50-byte record stride costs -32.21%, float-aware byte-plane slicing costs -31.60%. Adjacent triangles in an STL share vertices, so the raw stream carries local redundancy that LZ77 exploits and transposition destroys. On this data the fallback is already selecting correctly.

Declared gaps. Two domains present in the synthetic table have no real-world counterpart here. **LiDAR:** no source combining a suitable license with a usable format - USGS 3DEP proved unreachable and the one live alternative is LAZ, which would require laspy or PDAL and break the zero-dependency property of the reproducible script. **Audio:** no source offering direct WAV files at a proportionate size, the smallest real option costing 157 MB of transfer for a single usable file. A declared gap is preferable to a figure of uncertain provenance.

Domain Transform Specifications & Formal Theory

1. Astrophysics FITS & Multidimensional Tensor Modeling

Astronomical instruments (e.g., JWST NIRCam/MIRI, Hubble STIS) produce FITS data cubes formatted in Big-Endian standard byte order with 2880-byte header blocks followed by dense floating-point (>f4) or integer (>i4) matrices.

What the pipeline actually does today. A .fits input is classified by `DomainRouter.detect_domain()` as the generic `FLOAT32` domain and preconditioned by `ScientificFloatEngine.transform_f32_byte_plane()`. That transform is a single stage - a 4-plane byte transposition applied uniformly across the whole file, headers included:

$$\mathcal{M}_{\{i,j\}} \xrightarrow{\text{slicing}} \{P_0(i), P_1(i), P_2(i), P_3(i)\}$$

Floating-point matrices thereby isolate the sign bit and exponent (high byte `P_3`) from the least significant mantissa bits (`P_0`, `P_1`). Because physical sky backgrounds exhibit localized thermal equilibrium, `P_3` and `P_2` collapse into long run-length sequences, allowing entropy coders to process uniform statistical distributions. It is worth being precise about the limits of this path: it applies no spatial delta, performs no endianness handling, and has no awareness of FITS structure - it treats the file as an undifferentiated 4-byte-aligned buffer. The **+16.31%** reported in the table above is what this generic path achieves on a real JWST observation.

A dedicated engine exists but is not yet routed. The codebase also contains a `FitsEngine` class (`kolmox.engines.extended_domains`) implementing the richer two-stage scheme this section previously described as if it were the active path: a per-HDU horizontal modular delta,

$$\Delta_{\{x,y\}} = (S_{\{x,y\}} - S_{\{x-1,y\}}) \bmod{2^{32}}$$

followed by byte-plane slicing of the residuals, with `astropy`-based HDU parsing and its own serialization format. It is covered by unit tests, but `DomainRouter` never dispatches to it, and it does not emit a KMX2 container - so no figure in this paper is produced by it. Wiring it into the router, and measuring whether the delta stage actually improves on the generic transposition for astronomical data, is open work.

2. High-Throughput In-Memory Vector Slicing

In dense vector environments (e.g., neural embeddings, high-frequency telemetry caching), float representations incur severe penalties under LZ77-based match finders due to the chaotic nature of low-order mantissa bits. KolmoX transposes contiguous 32-bit buffers into 4 disjoint orthogonal memory blocks, isolating high-entropy IEEE-754 mantissa noise. Measured end-to-end through the pipeline, this path sustains **~642 MB/s** decompression on a 250k-value synthetic float buffer and **~714 MB/s** on the 117 MB JWST observation; the highest figure recorded across all domains is **~1094 MB/s**, on the fixed-stride binary register path.

KMX2 Container Specification

The **KMX2** container format implements a fixed 24-byte header supporting multi-stream encapsulation:

```
Header = < Magic(4B), Version(2B), DomainID(1B), Flags(1B), RawSize(8B), AuxLen(8B) >
```